Claude buys recall with volume, and credibility with precision

Across 10 investment-research seeds, a weaker model matches a stronger one's issue coverage by firing 9× more objections — half of them invented. The same failure appears in the plans themselves as arbitrary numeric thresholds.

So the metric that ranks discriminators is not recall or F1, but manufactured objections per review — and the eval must be built to make that visible.

Mayank Mahajan

Key takeaways

- **Recall does not rank the models, and F1 does not fix it.** Haiku 4.5 covers 53% of blocking gold issues — beating Sonnet 5 (29%) and Opus 4.8 (33%) — by raising 28.6 objections per review against Opus 5's 6.7. Only manufactured-per-review produces the expected gradient: 14.0 / 0.10 / 0.15 / 0.05.

	Haiku 4.5	Sonnet 5	Opus 4.8	Opus 5
Pairwise accuracy, strong items *(baseline 50%)*	58%	67%	67%	75%
Blocking-issue recall *(of 45)*	53%	29%	33%	62%
Manufactured objections / review	14.0	0.10	0.15	0.05

- **Confidence is a free difficulty label.** On items I marked `strong`, the ranking holds and all four beat chance. On `weak` items all four lose to always-answer-A — those items measure nothing. Score the sure items as the headline.
- **The same defect shows up in the plans.** *Manufactured precision* — invented numeric thresholds (a 20:1 prior, a 25× likelihood-ratio bar, 0.5%/0.2% entry gates) — appeared in 5 of 10 pairs. Unearned specificity reads as rigor.
- **Two generators are the whole ceiling.** Haiku ∪ Opus 5 covers 76% of blocking issues; Sonnet 5 and Opus 4.8 contribute zero unique findings.

The instrument: 13 tenets, grouped by when they apply

	Tenet
A · Frame	1 Understand the actual goal before judging anything · 2 Fit the plan to the firm and the person · 3 Establish the variant view — where is the edge?
B · Design	4 Simplest thing that could resolve the question, first · 5 Measure the actual target, keep the measure meaningful · 6 Match horizon to phenomenon, not data frequency · 7 Match evidence standard to what the setting supports · 8 Breadth over strength — test where it should *not* work
C · Execution	9 Sequence for information — get to 60 before going to 99 · 10 Assume AI execution; human checkpoints, calibrated explainability
D · Honesty	11 Always ask what would make this wrong · 12 Do not invent economics — stay inside the question as framed
E · Review	13 Standard objections must earn their place in *this* context

- Groups are **ordered** — a frame error (A) invalidates everything downstream — and precedence rules resolve the designed conflicts: rigor *vs* speed, breadth *vs* firm fit, explainability *vs* performance.
- **Tenet 13 is the precision tenet**, added because of the failure on slide 2. Every item carries firm context: a plan is good or bad **for someone**.
- Decisive tenets cluster hard — 2 and 7 drove five items each; 6, 9, 12 were never decisive.

Part III — spend the expert only where models are weakest

No generator proposes **24% of blocking issues**; recovering them needs an expert who has read both plans writing from scratch — irreducible at ~18 min. So **don't degrade the protocol, vary how many items receive it.**

#	Step	Expert min
1	Generate — Haiku 4.5 + Opus 5 critique both plans, all items (~70 candidates/item)	0
2	Adjudicate — Opus 5 buckets + dedups, confidence-flagged	0
3	Re-adjudicate — a second model forces a call on the *uncertain* pile only	0
4–5	Rank & split — score items by auto-tier weakness; top ~5 of 10 to full depth	0
6	Full-depth pass, 5 × 18 min — frame enums · read both plans · free-recall write, *candidates unseen*	90
7–9	Merge · audit ~4 auto-accepted cards · publish recall beside audited precision	2
	Total — 92 min / 10 items	9.2 / item

- **≈2× cheaper than free-hand (18 min/item) at F1 ≈ 0.90.** F1 0.95 costs 13.1 min.
- **Skip expert triage of candidates** — worst value in the protocol: 9.4 min for ~8pp.
- **The cost:** auto-accepting leaves Haiku's ~22 confidently-manufactured cards per item unlabelled — the very negatives that ranked the models.

Next steps

1. **Measure the one unmeasured number.** Adjudicator precision on candidates it confidently marks real is assumed, not known, and the whole Part III frontier rests on it. An audit of ~4 cards per 10 items settles it in the first session; at 0.70 rather than 0.95 the cost rises only to 11.5 min/item.
2. **Break the single-reviewer dependency.** All gold is one desk, so "*wrong*" and "*disagrees with me*" are nowhere separable. A second labeller on the same 10 items turns `weak` from a self-report into measured inter-expert disagreement.
3. **Fill the four context-flip items.** Same seed, same plans, different firm — the one axis isolating tenet 2. Measured gold-free only today, so context-sensitivity is unquantified.
4. **Rescore with a non-Opus-5 adjudicator.** Opus 5 scored its own reviews and posts the best recall *and* precision under its own judge; the Sonnet / Opus 4.8 inversion (97.0 *vs* 91.7) is where that could move the ranking.

Ordered by what each unblocks: (1) the cost model · (2) every accuracy number · (3) a missing axis · (4) the margin on one comparison.