

Claude buys recall with volume, and credibility with precision

Across 10 investment-research seeds, a weaker model matches a stronger one's issue coverage by firing 9× more objections — half of them invented. The same failure appears in the plans themselves as arbitrary numeric thresholds.

So the metric that ranks discriminators is not recall or F1, but **manufactured objections per review** — and the eval must be built to make that visible.

Mayank Mahajan

Key takeaways

- **Recall does not rank the models, and F1 does not fix it.** Haiku 4.5 covers 53% of blocking gold issues — beating Sonnet 5 (29%) and Opus 4.8 (33%) — by raising **28.6 objections per review against Opus 5's 6.7**. Only **manufactured-per-review** produces the expected gradient: **14.0 / 0.10 / 0.15 / 0.05**.

	Haiku 4.5	Sonnet 5	Opus 4.8	Opus 5
Pairwise accuracy, strong items (<i>baseline 50%</i>)	58%	67%	67%	75%
Blocking-issue recall (<i>of 45</i>)	53%	29%	33%	62%
Manufactured objections / review	14.0	0.10	0.15	0.05

- **Confidence is a free difficulty label.** On items I marked **strong**, the ranking holds and all four beat chance. On **weak** items **all four lose to always-answer-A** — those items measure nothing. Score the sure items as the headline.
- **The same defect shows up in the plans.** *Manufactured precision* — invented numeric thresholds (a 20:1 prior, a 25× likelihood-ratio bar, 0.5%/0.2% entry gates) — appeared in **5 of 10** pairs. Unearned specificity reads as rigor.
- **Two generators are the whole ceiling.** Haiku ∪ Opus 5 covers **76%** of blocking issues; **Sonnet 5 and Opus 4.8 contribute zero unique findings**.

The instrument: 13 tenets, grouped by when they apply

	Tenet
A · Establish the frame	1 Understand the actual goal before judging anything · 2 Fit the plan to the firm and the person · 3 Establish the variant view — where is the edge?
B · Is the design right for that frame	4 Simplest thing that could resolve the question, first · 5 Measure the actual target, keep the measure meaningful · 6 Match horizon to phenomenon, not data frequency · 7 Match evidence standard to what the setting supports · 8 Breadth over strength — test where it should <i>not</i> work
C · Execution	9 Sequence for information — get to 60 before going to 99 · 10 Assume AI execution; human checkpoints, calibrated explainability
D · Intellectual honesty	11 Always ask what would make this wrong · 12 Do not invent economics — stay inside the question as framed
E · The review itself	13 Standard objections must earn their place in <i>this</i> context

- Groups are **ordered**: a frame error (A) invalidates everything downstream, so precedence rules resolve the designed conflicts — rigor *vs* speed, breadth *vs* firm fit, explainability *vs* performance.
- Tenet 13 is the precision tenet** and exists because of the failure on slide 2. Every item carries firm context,

Part III — spend the expert only where models are weakest

Models cannot produce the 24% of blocking issues no generator proposes; recovering those needs an expert who has read both plans writing from scratch. That pass is irreducible at ~18 min. So don't degrade the protocol — vary how many items receive it.

#	Step	Expert min
1	Generate — Haiku 4.5 + Opus 5 critique both plans, every item (~70 candidates/item)	0
2	Adjudicate — Opus 5 buckets + dedups, confidence-flagged	0
3	Re-adjudicate — a second model forces a call on the uncertain pile only	0
4-5	Rank & split — score items by auto-tier weakness; top ~5 of 10 to full depth	0
6	Full-depth pass, 5 items × 18 min — frame enums · read both plans · free-recall write, candidates unseen	90
7-9	Merge · audit ~4 auto-accepted cards · publish recall beside audited precision	2
	Total — 92 min / 10 items	9.2 / item

Next steps

1. **Measure the one unmeasured number.** Adjudicator precision on candidates it confidently marks real is assumed, not known — the whole Part III frontier rests on it. An audit of ~4 cards per 10 items settles it in the first session; at 0.70 rather than 0.95 the cost rises only to 11.5 min/item.
2. **Break the single-reviewer dependency.** All gold is one desk, so "*wrong*" and "*disagrees with me*" are not separable anywhere in this repo. A second labeller on the same 10 items converts **weak** from a self-report into measured inter-expert disagreement.
3. **Fill the four context-flip items.** Same seed, same plans, different firm — the one axis that isolates tenet 2. Currently measured gold-free only, so context-sensitivity is unquantified.
4. **Rescore with a non-Opus-5 adjudicator.** Opus 5 scored its own reviews and posts the best recall *and* precision under its own judge; the Sonnet / Opus 4.8 inversion (97.0 *vs* 91.7) is where that could be moving the ranking.

Ordered by what each unblocks: (1) the cost model · (2) every accuracy number · (3) a missing axis · (4) the margin on one comparison.